



Live full-pipeline scaling; finite workload, shared serving telemetry

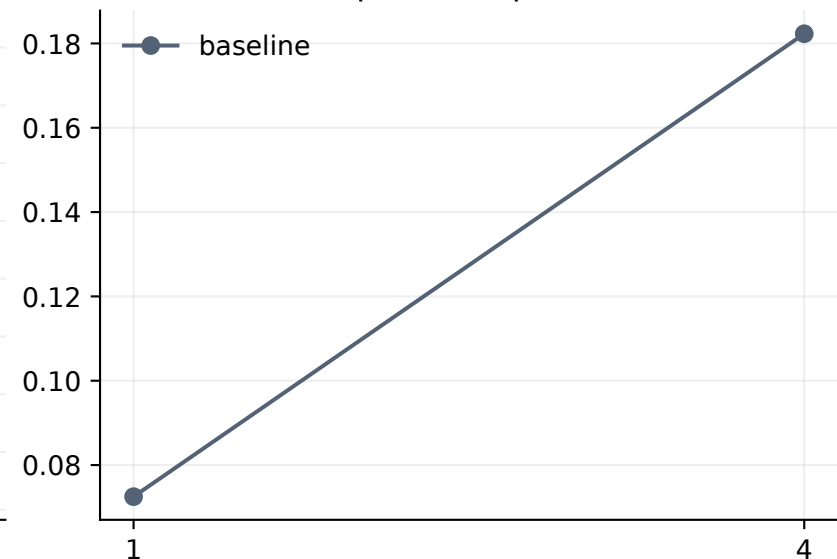
Mean full response (s)



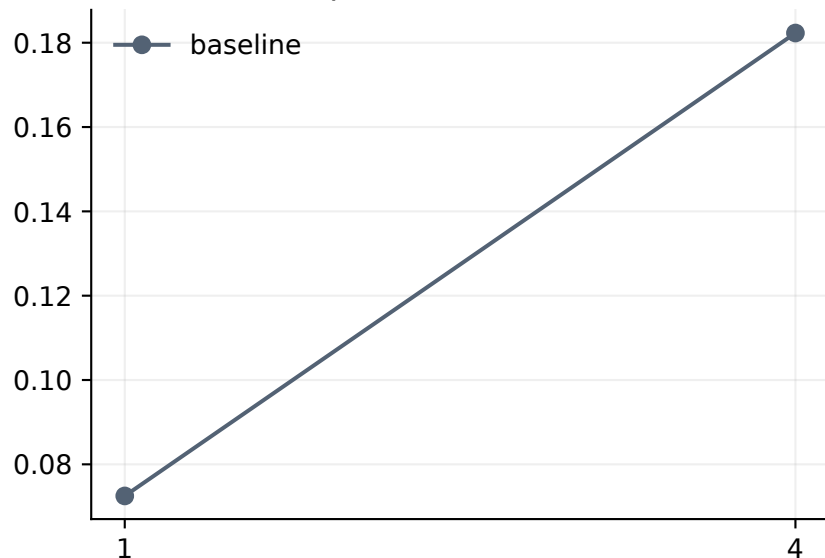
P95 full response (s)



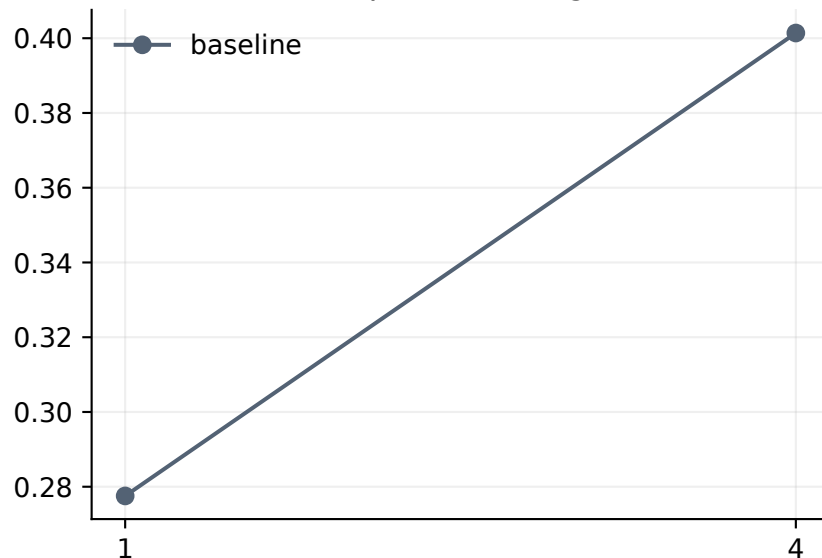
Completed requests / s



Completed within 30 s / s



Orchestrator peak KV usage (fraction)



Orchestrator peak waiting requests

